
DataSci 400

lesson 4: dealing with numeric data

Seth Mottaghinejad

today's agenda

- types of numeric columns
- two types of functions
 - **summarization** function
 - **transformation** functions
- common transformations for **feature engineering**:
 - binning, trimming, normalizing
- special transformations
 - RFM (recency, frequency, monetary)

types of numeric columns

- **floats:** mathematicians call it a **continuous variable**
 - in `pandas` we use `float64` (or `float32` etc.)
- **integers:** mathematicians call it a **discrete variable**
 - in `pandas` we use `int64` (or `int32` etc.)
- **datetimes:** this is more like a construct, behind the scenes datetimes are like integer columns with the counter starting at some date in the past like `1960-01-01 00:00:00`
 - in `pandas` we use `datetime64`
 - the **difference** between two `datetime64` is a `timedelta[n]` type

summary vs transformation functions

- a **summary** function takes in **an array of numbers** and returns a **single number** (or several sometimes)
 - **average:** `mean([4, 5, 6])` → `5`
- a **transformation** function takes in **an array of numbers** and returns **an array of numbers of the same size** where we assume a 1-1 correspondence between elements
 - **log:** `log([4, 5, 6])` → `[1.386, 1.610, 1.792]`
 - **round:** `log([1.386, 1.610, 1.792])` → `[1.4, 1.7, 1.8]`
 - **lag:** `log([4, 5, 6])` → `[NA, 4, 5]`

types of numeric columns

- feature engineering is all about **applying transformation** functions to the data to create **new features from existing features**
- for numeric features three common transformations are
 - **binning**: take a numeric feature and return a categorical feature based on what **pre-defined interval** the number is in
 - **trimming**: remove outliers from the data by replacing values that exceed a threshold with the threshold or NA
 - **normalization**: rescale columns in the data so they all have the same **scale** (range of values)

two common ways to normalize

- **Z-normalization** makes most values fall between -2 and 2

$$x \rightarrow \frac{x - \text{mean}(x)}{\text{std}(x)}$$

- **min-max normalization** forces all the values between 0 and 1

$$x \rightarrow \frac{x - \min(x)}{\max(x) - \min(x)}$$

- **some (but not all)** ML algorithms require normalized data

special transformations: RFM example

- because datasets can be very different depending on industry, there are some data transformations that are **industry-specific**
- in **retail** for example, where we have **times series** data of customer purchases we can extract RFM features
 - **R is for recency**: how long ago was the most recent purchase?
 - **F is for frequency**: how many purchases were made in the last week (or month, or whatever window you choose)?
 - **M is for monetary**: how much in total was spent in the last week (or month, or whatever window you choose)?

date	customer	purchased
2020-01-01	XYZ	\$50
2020-01-03	XYZ	\$23
2020-02-11	XYZ	\$35
2020-01-02	ABC	\$50
2020-01-02	ABC	\$23
2020-01-29	ABC	\$35
⋮	⋮	⋮

date	customer	purchased	R	F	M
2020-01-01	XYZ	\$50	NA	NA	NA
2020-01-03	XYZ	\$23	2 days	2	\$73
2020-01-11	XYZ	\$35	8 days	1	\$35
2020-01-02	ABC	\$65	NA	NA	NA
2020-01-03	ABC	\$25	1 day	2	\$90
2020-01-29	ABC	\$35	26 days	1	\$35
⋮	⋮	⋮	⋮	⋮	⋮

the end